

[Home](#)[First Task Quick Start](#)[FAQ](#)

LEARN

[Tasking hub](#)[The Rubric](#)[Foundations](#)[Knowledge Graph](#)[Image skills](#)[How to see](#)[Visual quality](#)[Find realism & Artifacts](#)[AI artifact](#)[How to mark severity](#)[Audio skills](#)[Figure out sound vs speech](#)[Caption environmental sounds](#)[Caption speech and lyrics](#)[Transcribe lyrics correctly](#)[Pacing and emotion](#)[Audio artifacts](#)[Audio coherence](#)[Video skills](#)[Segmentation skills](#)

IMAGE - EVALS

[Overview](#)[Spotting subtle failures](#)[H2H Image FAQs](#)[T2I text-to-image](#)[I2I image-to-image](#)[UD Caption Elo](#)[Omni R2I Elo](#)

IMAGE - NON-EVALS

[Overview](#)[Annotation Critic Full - Self](#)

How to mark severity

Your vote strength should reflect **how big the difference between the two images really is** — not how confident you feel. This guide shows when to choose Tie, Slightly, or Strongly, with two worked comparisons.

The rating scale

You'll use this five-point scale for **Overall Preference** and for **each dimension** you score — instruction following, reference preservation, visual quality, and AI artifacts. It works the same whether the task is **omni-r2i-elo**, **text-to-image-compare**, or **omni-elo**.

Strongly Prefer A

A is clearly and significantly stronger.

Slightly Prefer A

A has a noticeable but small edge.

Tie

Truly equal — rare.

Slightly Prefer B

B has a noticeable but small edge.

Strongly Prefer B

B is clearly and significantly stronger.

Slightly vs. Strongly — how to tell them apart

Slightly prefer

The two images are broadly comparable and the call takes some deliberation. Usually one of:

- you can name **the one detail** that sets one response apart from the other
- both look similar and the only difference **shows up on a close look**, or
- another rater could **reasonably land the other way**

Example: on AI artifacts, Response A has 3 visible artifacts and Response B has 2 — just a one-artifact gap → **Slightly Prefer B**.

Strongly prefer

The winner is **obvious on the first look**. Usually one of:

- the losing image has a **jarring failure a casual viewer would immediately flag** (bad anatomy, a required element missing or wrong, the reference style broken, an obvious artifact), or
- you **wouldn't second-guess it**, and you wouldn't expect another rater to land differently

H

Home

First Task Quick Start

FAQ

LEARN

Tasking hub

The Rubric

Foundations

Knowledge Graph

Image skills

How to see

Visual quality

Find realism & Artifacts

AI artifact

How to mark severity

Audio skills

Figure out sound vs
speechCaption environmental
soundsCaption speech and
lyricsTranscribe lyrics
correctly

Pacing and emotion

Audio artifacts

Audio coherence

Video skills

Segmentation skills

IMAGE - EVALS

Overview

Spotting subtle failures

H2H Image FAQs

T2I text-to-image

I2I image-to-image

UD Caption Elo

Omni R2I Elo

IMAGE - NON-EVALS

Overview

Annotation Critic Full - Self

Example: Response A has one clear problem you'd spot right away (say, a badly deformed hand) while Response B looks clean — or has only a tiny flaw you'd have to hunt for → **Strongly Prefer B.**

Tie

Neither image is meaningfully better after a careful look. **Truly rare.**

Example: both responses have the same 2 artifacts at the same severity and nothing else separates them → **Tie** (still name what you compared and why it came out even).

The idea is the same in every dimension: a small difference in count or severity is a *Slightly*, a large or obvious difference is a *Strongly*, and a genuinely even split is a *Tie*.

Rule of thumb

Imagine someone glanced at both images. *Slightly* — they'd pick the winner when you ask them to choose. *Strongly* — they'd point out the problem unprompted.

A "Strongly" winner can still have its own small flaws — what matters is whether the *gap* between the two is obvious. And a tie on one dimension doesn't mean the images are tied overall.

Example 1 — a *Slightly* call

EDITING PROMPT

"add the protective headgear a construction worker wears on site, change nothing else"



Reference

H

- Home
- First Task Quick Start
- FAQ

LEARN

- Tasking hub
- The Rubric
- Foundations
 - Knowledge Graph
 - Image skills
 - How to see
 - Visual quality
 - Find realism & Artifacts
 - AI artifact
 - How to mark severity
 - Audio skills
 - Figure out sound vs speech
 - Caption environmental sounds
 - Caption speech and lyrics
 - Transcribe lyrics correctly
 - Pacing and emotion
 - Audio artifacts
 - Audio coherence
 - Video skills
 - Segmentation skills

IMAGE - EVALS

- Overview
- Spotting subtle failures
- H2H Image FAQs
- T2I text-to-image
- I2I image-to-image
- UD Caption Elo
- Omni R2I Elo

IMAGE - NON-EVALS

- Overview
- Annotation Critic Full - Self



Output A



Output B

WHAT YOU SEE

Both outputs add the hard hat cleanly. Both also subtly flatten the man's head shape, which the prompt didn't ask for — so that flaw appears on *both* sides and doesn't separate them. The real difference: **Output B adds three unnecessary white lines on his forehead; Output A doesn't.**


Overall Preference — Slightly Prefer A

Both images follow the core instruction and are fully usable. The only real separator is B's small unrequested detail (the white lines), which you notice on a close look. The shared head-shape change doesn't decide the call — so the gap is real but minor, making it a Slightly, not a Strongly.

Example 2 — a *Strongly* call

EDITING PROMPT

"Add a large, hyper-realistic human eyeball resting on the wooden floor in the immediate foreground of the hallway, centered on the floorboards and facing upward toward the ceiling, as if it has fallen from the artwork at the end."



Home

First Task Quick Start

FAQ

LEARN

Tasking hub

The Rubric

Foundations

Knowledge Graph

Image skills

How to see

Visual quality

Find realism & Artifacts

AI artifact

How to mark severity

Audio skills

Figure out sound vs speech

Caption environmental sounds

Caption speech and lyrics

Transcribe lyrics correctly

Pacing and emotion

Audio artifacts

Audio coherence

Video skills

Segmentation skills

IMAGE - EVALS

Overview

Spotting subtle failures

H2H Image FAQs

T2I text-to-image

I2I image-to-image

UD Caption Elo

Omni R2I Elo


IMAGE - NON-EVALS

Overview

Annotation Critic Full - Self



Output A



Output B

WHAT YOU SEE

Output A adds the eyeball exactly as described — large, hyper-realistic, on the floorboards — and preserves the reference's style and everything else. **Output B recolors the eyeball and shifts the artistic style**, so it looks out of place and creepy rather than blending in, and it even colorizes the man's eye for no reason — changes nobody asked for.

Overall Preference — Strongly Prefer A

B breaks the reference's style and adds several unrequested changes that are obvious at a glance, while A follows the prompt cleanly and preserves the source. The prompt-writer would spot the difference immediately — a big, clear gap, so it's a Strongly.

The most common mistake: under-calling

Fellows tend to default to *Slightly* to play it safe. Before you settle for a Slightly, check the three cases below — each one is a Strongly.

You can call it a Strongly when

- All the differentiating flaws sit on one side — one image is clean, the other carries the problems.
- One image misses something the prompt required — a missing or wrong required element is a Strongly, even if the rest looks similar.
- The same flaw repeats — the same small artifact appearing 4 times vs. once is a Strongly, not a Slightly. Count matters.

Bottom line: don't shrink your vote just because you're not 100% sure. Size the vote to the gap between the images, not to your confidence.



- Home
- First Task Quick Start
- FAQ
- LEARN
 - Tasking hub
 - The Rubric
 - Foundations
 - Knowledge Graph
 - Image skills
 - How to see
 - Visual quality
 - Find realism & Artifacts
 - AI artifact
 - How to mark severity
 - Audio skills
 - Figure out sound vs speech
 - Caption environmental sounds
 - Caption speech and lyrics
 - Transcribe lyrics correctly
 - Pacing and emotion
 - Audio artifacts
 - Audio coherence
 - Video skills
 - Segmentation skills
- IMAGE - EVALS
 - Overview
 - Spotting subtle failures
 - H2H Image FAQs
 - T2I text-to-image
 - I2I image-to-image
 - UD Caption Elo
 - Omni R2I Elo
- IMAGE - NON-EVALS
 - Overview
 - Annotation Critic Full - Self

Knowledge checks

KNOWLEDGE CHECK

The prompt asked to add a red umbrella to the person's hand. Response A adds it cleanly and everything else looks natural. Response B adds the umbrella too, but the person's hand is clearly deformed. Nothing else separates the two. What's your Overall Preference?

Strongly Prefer A

Slightly Prefer A

Tie

Slightly Prefer B

Strongly Prefer B

Check answer

KNOWLEDGE CHECK

Both responses follow the prompt and look good overall. Only on a close look do you notice Response A has 3 minor artifacts (a couple of soft, slightly-off edges) and Response B has 2 of the same kind. Nothing else separates them. What's your Overall Preference?

Strongly Prefer A

Slightly Prefer A

Tie

Slightly Prefer B

Strongly Prefer B

Check answer

Finished every knowledge check? Submit to record your completion.

Submit

Previous

Next